

XNU: Reference Count Leak in POSIX Shared Memory Object

Zhuo Liang

April 2, 2019

Background

XNU supports **Shared Memory**¹ for inter process communication. The kernel provides two kinds of memory-sharing mechanisms: POSIX shared memory and System V shared memory.

```
1 int main( int argc, char** argv ) {
2     int fd;
3     unsigned* addr;
4     /* Create a new memory object */
5     fd = shm_open( "/bolts", O_RDWR | O_CREAT, 0777 );
6     /* Set the memory object's size */
7     ftruncate( fd, sizeof( *addr ) );
8     /* Map the memory object */
9     addr = mmap( 0, sizeof( *addr ), PROT_READ | PROT_WRITE,
10                 MAP_SHARED, fd, 0 );
11
12     /* Write to shared memory */
13     *addr = 1;
14
15     /* The memory object remains in the system after the close */
16     close( fd );
17     /*
18      * To remove a memory object you must unlink it like a file.
19      * This may be done by another process.
20      */
21     shm_unlink( "/bolts" );
22     return EXIT_SUCCESS;
23 }
```

Listing 1: POSIX shared memory usage

¹https://en.wikipedia.org/wiki/Shared_memory

Listing 1 is a [typical usage](#)² of POSIX shared memory. The steps can be divided into following items:

1. **shm_open** Create a new shared memory object and put it into cache. Several times of **shm_open** is supported and they will share the common shared memory object from kernel cache.
2. **ftruncate** Allocate backend sharing memory for shared memory object and this operation will mark the object as **PSHM_ALLOCATED**.
3. **mmap** Map the allocated memory into the process's space and the returned value is the start of the shared memory.
4. **Direct Read/Write** Since the memory is already mapped in the process's task space, the process can read and write the shared memory now.
5. **close** Release the file descriptor and decrease the reference of the shared memory object.
6. **shm_unlink** Unlink the path, this operation would decrease the reference count of shared memory object and mark the object as **PSHM_REMOVED**.

Leak Issue

This issue is about the management of shared memory object. The **close** operation of POSIX shared memory object is **pshm_closefile** which will call **pshm_close**.

```
1 // bsd/kern/posix_shm.c
2 static int
3 pshm_close(struct pshminfo * pinfo, int dropref)
4 {
5     int error = 0;
6     struct pshmobj *pshmobj, *pshmobj_next;
7
8     /*
9      * If we are dropping the reference we took on
10     * the cache object, don't enforce the
11     * allocation requirement.
```

²[www.qnx.com shared memory](http://www.qnx.com/shared_memory)

```

12     */
13     if (!dropref &&
14         ((pinfo->pshm_flags & PSHM_ALLOCATED)
15          != PSHM_ALLOCATED)) { // [a]
16         return (EINVAL);
17     }
18     /* DIAGNOSTIC */
19     pinfo->pshm_usecount--; /* release this fd's reference */ // [b]
20     ...

```

Listing 2: `pshm_close` function

At **[a]**, `PSHM_ALLOCATED` is checked and this flag is only set in `ftruncate`. `pshm_closefile` passes 0 as the second parameter to this function and this means if we open the shared memory and close at once, the `pshm_usecount` will not be decreased. Let's see what would happen if we do following steps:

```

1     const char *shm_name = "/test.shm";
2     int shm_fd = shm_open(shm_name, O_RDWR | O_CREAT, 0666); // [c]
3 #define MAX_OPEN_TIMES 0xff
4     for (size_t i = 0; i < MAX_OPEN_TIMES; i++) {
5         int reopen_shm_fd = shm_open(shm_name, O_RDWR); // [d]
6         close(reopen_shm_fd); // [e]
7     }

```

Listing 3: Corrupting steps

1. **[c]** Create the memory object and the put it into cache, the `pshm_usecount` is 2 now. One for file descriptor and the other for cache.
2. **[d]** Open the same path, this will search the object from kernel cache and increase the `pshm_usecount`.

```

1 // bsd/kern/posix_shm.c
2 int
3 shm_open(proc_t p, struct shm_open_args *uap, int32_t *retval) {
4     /*
5      * If we find the entry in the cache, this
6      * will take a reference, allowing us to
7      * unlock it for the permissions check.
8      */
9     error = pshm_cache_search(&pinfo, &nd, &pcache, 1);
10 }

```

Listing 4: Search object from cache

3. **[e]** Close the reopened file descriptor, recall the aforementioned close operation that this will not decrease the `pshm_usecount` because `ftruncate` has not been called yet.

After above steps, the `pshm_usecount` will be `0x101`. One for `shm_fd` which we still hold, one for cache and `0xff` for reopened file descriptors which we already closed. And the `pshm_usecount` is a 32 bit integer which means if we set `MAX_OPEN_TIMES` to `0xffffffff`, the result of `pshm_usecount` will be 1, but we still hold one file descriptor. If we `unlink` the path, which will decrease the usecount and of course release the object memory, and then do anything on that file descriptor, an use-after-free issue occurs!

Fixing

Apple adds a function named `pshm_deref`, which will be called when closing a handle of POSIX shared memory object or unlinking the path, to fix this issue.

```
1 int64 pshm_deref(__int64 a1, __int64 a2)
2 {
3     __int64 v2; // rsi
4     int v3; // eax
5     __int64 result; // rax
6     _QWORD *i; // rbx
7     _DWORD *v6; // r13
8     __int64 v7; // rax
9     __int64 v8; // [rsp-8h] [rbp-30h]
10
11     v8 = a1;
12     v2 = 1LL;
13     lck_mtx_assert(&psx_shm_subsys_mutex, 1LL);
14     v3 = *(_DWORD *) (a2 + 4);
15     result = (unsigned int) (v3 - 1);
16     *(_DWORD *) (a2 + 4) = result;
17     if ( !(_DWORD) result )
18     {
19         ...
20     }
21     return result;
22 }
```

Listing 5: `pshm_deref` function

Timeline

1. **2018/12/04** Discovery of this issue.
2. **2018/12/11** Reported to `product-security@apple.com`.
3. **2019/03/13** Checked that the issue was fixed in Beta4.